
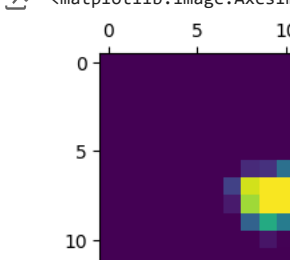



Problem Statement 2: To Implement Feed forward Neural Network with keras and tensorflow. Import the necessary packages Load the training and testing data(MNIST) Define the network architecture using keras Train the model using SGD Evaluate the network Plot the training loss and accuracy

 `<matplotlib.image.AxesImage at 0x7ab9a8dad4e0>`



A 25x25 pixel plot showing a yellow handwritten digit '3' on a dark purple background. The plot has x and y axes labeled from 0 to 25. The digit is composed of yellow and light green pixels, with some darker green pixels at the edges. The background is a solid dark purple.

```
 array([[0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ],  
         [0.,      , 0.,      , 0.,      , 0.,      , 0.,      ]])
```

```

0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      ],
[0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.01176471, 0.07058824, 0.07058824,
0.07058824, 0.49411765, 0.53333333, 0.68627451, 0.10196078,
0.65098039, 1.      , 0.96862745, 0.49803922, 0.      ,
0.      , 0.      , 0.      ],
[0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.11764706, 0.14117647,
0.36862745, 0.60392157, 0.66666667, 0.99215686, 0.99215686,
0.99215686, 0.99215686, 0.99215686, 0.88235294, 0.6745098 ,
0.99215686, 0.94901961, 0.76470588, 0.25098039, 0.      ,
0.      , 0.      , 0.      ],
[0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.19215686, 0.93333333, 0.99215686,
0.99215686, 0.99215686, 0.99215686, 0.99215686, 0.99215686,
0.99215686, 0.99215686, 0.98431373, 0.36470588, 0.32156863,
0.32156863, 0.21960784, 0.15294118, 0.      , 0.      ,
0.      , 0.      , 0.      ],
[0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.07058824, 0.85882353, 0.99215686,
0.99215686, 0.99215686, 0.99215686, 0.99215686, 0.77647059,
0.71372549, 0.96862745, 0.94509804, 0.      , 0.      ,
0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      ],
[0.      , 0.      , 0.      , 0.      , 0.      ,
0.      , 0.      , 0.      , 0.31372549, 0.61176471,
0.41960784, 0.99215686, 0.99215686, 0.80392157, 0.04313725,
0.      , 0.16862745, 0.60392157, 0.      , 0.      ,

```

```

model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28,28)),
    keras.layers.Dense(128, activation = 'relu'),
    keras.layers.Dense(10, activation='softmax')
])

```

→ /usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_c`
super().__init__(**kwargs)

```
model.summary()
```

→ Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 128)	100,480
dense_4 (Dense)	(None, 10)	1,290

Total params: 101,770 (397.54 KB)
Trainable params: 101,770 (397.54 KB)
Non-trainable params: 0 (0.00 B)

```

model.compile(
    optimizer='sgd',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])

```

```
history = model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=10)
```

→ Epoch 1/10
1875/1875 ————— 6s 3ms/step - accuracy: 0.7348 - loss: 1.0318 - val_accuracy: 0.9040 - val_loss: 0.3550
Epoch 2/10
1875/1875 ————— 6s 3ms/step - accuracy: 0.9011 - loss: 0.3529 - val_accuracy: 0.9196 - val_loss: 0.2932
Epoch 3/10
1875/1875 ————— 5s 3ms/step - accuracy: 0.9157 - loss: 0.3000 - val_accuracy: 0.9264 - val_loss: 0.2627
Epoch 4/10
1875/1875 ————— 6s 3ms/step - accuracy: 0.9236 - loss: 0.2731 - val_accuracy: 0.9319 - val_loss: 0.2420
Epoch 5/10
1875/1875 ————— 10s 3ms/step - accuracy: 0.9315 - loss: 0.2419 - val_accuracy: 0.9370 - val_loss: 0.2235
Epoch 6/10
1875/1875 ————— 6s 3ms/step - accuracy: 0.9375 - loss: 0.2218 - val_accuracy: 0.9415 - val_loss: 0.2053
Epoch 7/10
1875/1875 ————— 11s 4ms/step - accuracy: 0.9424 - loss: 0.2030 - val_accuracy: 0.9447 - val_loss: 0.1937
Epoch 8/10
1875/1875 ————— 10s 3ms/step - accuracy: 0.9453 - loss: 0.1895 - val_accuracy: 0.9482 - val_loss: 0.1813
Epoch 9/10
1875/1875 ————— 6s 3ms/step - accuracy: 0.9493 - loss: 0.1787 - val_accuracy: 0.9521 - val_loss: 0.1705
Epoch 10/10
1875/1875 ————— 11s 4ms/step - accuracy: 0.9538 - loss: 0.1672 - val_accuracy: 0.9545 - val_loss: 0.1630

```
test_loss, test_acc=model.evaluate(x_test,y_test)
```

```
print("Loss-%.3f" %test_loss)
```

```
print("Accuracy=%.3f" %test_acc)
```

313/313 — 1s 2ms/step - accuracy: 0.9455 - loss: 0.1924
Loss=0.163
Accuracy=0.955

```
n=random.randint(0,9999)
```


```
plt.imshow(x_test[n])
```

```
plt.show
```

matplotlib.pyplot.show
def show(*args, **kwargs)

****Auto-show in jupyter notebooks****

The jupyter backends (activated via ``%matplotlib inline``, ``%matplotlib notebook``, or ``%matplotlib widget``), call ``show()`` at the end of every cell by default. Thus, you usually don't have to call it explicitly there.



```
import matplotlib.pyplot as plt
```

```
# Function to plot the training and validation accuracy/loss
```

```
def plot_accuracy_loss(history):
```

```
    # Set figure size
```

```
    plt.figure(figsize=(12, 5))
```

```
    # Plot training & validation accuracy values
```

```
    plt.subplot(1, 2, 1)
```

```
    plt.plot(history.history['accuracy'], label='Train Accuracy')
```

```
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
    plt.title('Model Accuracy')
```

```
    plt.ylabel('Accuracy')
```

```
    plt.xlabel('Epoch')
```

```
    plt.legend(loc='upper left')
```

```
    # Plot training & validation loss values
```

```
    plt.subplot(1, 2, 2)
```

```
    plt.plot(history.history['loss'], label='Train Loss')
```

```
    plt.plot(history.history['val_loss'], label='Validation Loss')
```

```
    plt.title('Model Loss')
```

```
    plt.ylabel('Loss')
```

```
    plt.xlabel('Epoch')
```

```
    plt.legend(loc='upper left')
```


```
    # Show the plot
```

```
    plt.tight_layout()
```

```
    plt.show()
```

```
# Assuming 'history' object is available from model.fit()
```

```
plot_accuracy_loss(history)
```



```
test_predict = model.predict(x_test)

# Get the classification labels

test_predict_labels = np.argmax(test_predict, axis=1)

confusion_matrix = tf.math.confusion_matrix(labels=y_test, predictions=test_predict_labels)

print("Confusion matrix of the test set:\n", confusion_matrix)
```

313/313 ————— 1s 2ms/step

Confusion matrix of the test set:

```
tf.Tensor(
[[ 966   0    1    1    0    4    6    1    1    0]
 [   0 1117    2    2    0    1    3    2    8    0]
 [   6    2  984   10    4    2    5    7    8    4]
 [   0    0    7  968    1   10    0    9   11    4]
 [   1    0    6    0  940    0    7    3    4   21]
 [   8    1    0   19    4  832   10    2   10    6]
 [  10    3    2    1    7    9  922    1    3    0]
 [   0    8   18    5    4    2    0  969    2   20]
 [   7    2    2   17    7   10    9    8  905    7]
 [   9    7    1   10   25    5    1    7    2  942]], shape=(10, 10), dtype=int32)
```

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.